

Central Reasoning Engine (CRE)


— Detailed Design

Version: 0.1 (First Iteration) **Status:** Draft **Specification:** See [CENTRAL_REASONING_ENGINE.md](#) **Depends on:** h-M3 v2, h-DL Context Vault v3

1. Architectural Overview

1.1 Placement in Puffin


The CRE is a core module in Puffin's main process. It sits between the orchestration layer (which manages the user workflow) and the 3CLI execution layer.



1.2 Module Structure

```
src/main/cre/
├── index.js          # Public API, IPC registration
├── plan-generator.js # Plan generation logic
├── ris-generator.js  # RIS generation logic
├── code-model.js     # h-DL instance management
├── schema-manager.js # h-DL schema evolution
├── introspector.js   # Post-implementation code analysis
└── assertion-generator.js # Inspection assertion generation
  └── lib/
    ├── hdsł-types.js  # h-M3 primitive type definitions
    ├── hdsł-validator.js # Schema/instance validation
    └── context-builder.js # Navigation context for code model queries
      └── ris-formatter.js # RIS markdown formatting
```

1.3 Component Diagram



2. Component Design

2.1 CRE Entry Point (`index.js`)

The entry point initializes all sub-components and registers IPC handlers.

```
/**  
 * @module CRE  
 * Central Reasoning Engine - core Puffin module.  
 * Provides plan generation, RIS generation, and code model maintenance.  
 */  
module.exports = {  
  /**  
   * Initialize the CRE with application context.  
   * @param {Object} context - { ipcMain, app, db, config, projectRoot }  
   */  
  async initialize(context) { /* ... */ },  
  
  /**  
   * Graceful shutdown. Persists any pending state.  
   */  
  async shutdown() { /* ... */ }  
};
```

IPC Channels Registered:


Channel	Direction	Purpose
	invoke	Start plan generation for a sprint

Channel	Direction	Purpose
cre:generate-plan		
cre:refine-plan	invoke	Submit user feedback for plan refinement
cre:approve-plan	invoke	Approve a plan, triggering assertion generation
cre:generate-ris	invoke	Generate RIS for a user story
cre:update-model	invoke	Trigger code model update after implementation
cre:query-model	invoke	Query the code model (for UI or debugging)
cre:get-plan	invoke	Retrieve a plan by sprint ID
cre:get-ris	invoke	Retrieve RIS by user story ID

2.2 PlanGenerator (`plan-generator.js`)

Responsible for FR-01 through FR-06.

State Machine



Key Methods

```
class PlanGenerator {
    /**
     * Analyze sprint and user stories, identify ambiguities.
     * @param {string} sprintId
     * @returns {Object} { questions: string[], draft: Plan | null }
     */
    async analyzeSprint(sprintId) { /* ... */ }

    /**
     * Generate implementation plan using code model context.
     * @param {string} sprintId
     * @param {Object} answers - Answers to clarifying questions
     * @returns {Plan}
     */
    async generatePlan(sprintId, answers) { /* ... */ }

    /**
     * Refine plan based on user feedback.
     * @param {string} planId
     * @param {string} feedback - User's change request
     * @returns {Plan}
     */
    async refinePlan(planId, feedback) { /* ... */ }

    /**
     * Approve plan and generate inspection assertions.
     * @param {string} planId
     * @returns {Plan} - Plan with assertions populated
     */
    async approvePlan(planId) { /* ... */ }
}
```

Plan Data Structure

```
/**
 * @typedef {Object} Plan
 * @property {string} id - UUID
 * @property {string} sprintId - FK to sprint
 * @property {string} status - idle|analyzing|questions_pending|generating|review_pending|approved
 * @property {PlanItem[]} items - Ordered implementation steps
 * @property {string[]} pendingQuestions - Unanswered clarifying questions
 * @property {Object} metadata - { createdAt, updatedAt, version }
 */

/**
 * @typedef {Object} PlanItem
 * @property {number} order - Execution sequence
 * @property {string} userStoryId - FK to user story
 * @property {string} branch - Git branch name
 * @property {string[]} dependsOn - IDs of prerequisite plan items
 * @property {string} summary - What this step accomplishes
 * @property {string} approach - How it will be implemented
 * @property {InspectionAssertion[]} assertions - Post-implementation checks
 */

/**
 * @typedef {Object} InspectionAssertion
 * @property {string} type - file_exists|function_exists|export_exists|pattern_match
 * @property {string} target - File path, function name, or pattern
 * @property {string} description - Human-readable description
 */
```

2.3 RISGenerator (`ris-generator.js`)

Responsible for FR-07 through FR-11.

Key Methods

```
class RISGenerator {  
    /**  
     * Generate RIS for a user story.  
     * Consults code model and approved plan.  
     * @param {string} userStoryId  
     * @returns {RIS[]}  
     */  
    async generateRIS(userStoryId) { /* ... */ }  
  
    /**  
     * Query code model for artifacts relevant to this story.  
     * Uses h-DSL navigation (PEEK, FOCUS, TRACE) to build context.  
     * @param {string} userStoryId  
     * @param {PlanItem} planItem  
     * @returns {Object} relevantContext  
     */  
    async consultModel(userStoryId, planItem) { /* ... */ }  
}
```

RIS Data Structure

```
/**  
 * @typedef {Object} RIS  
 * @property {string} id - UUID  
 * @property {string} userStoryId - FK to user story  
 * @property {string} planId - FK to plan  
 * @property {string} sprintId - FK to sprint  
 * @property {string} branch - Target git branch  
 * @property {string} content - Markdown content (the actual specification)  
 * @property {string} status - generated|sent|completed  
 * @property {Object} metadata - { createdAt, codeModelVersion }  
 */
```

RIS Content Format

Each RIS is a markdown file with the following structure:

```
# RIS: [User Story Title]  
  
## Context  
- Branch: `feature/xyz`  
- Depends on: [list of prerequisite branches/RIS]  
- Code Model Version: [hash]  
  
## Objective  
[One-sentence statement of what this implementation achieves]  
  
## Instructions  
  
### 1. [First task]  
- File: `src/path/to/file.js`  
- Action: create | modify | delete  
- Details:  
  [Specific implementation instructions]  
  
### 2. [Second task]  
...  
  
## Conventions  
- [Relevant patterns from the code model]  
- [Naming conventions, error handling patterns, etc.]  
  
## Assertions  
- [ ] [Inspection assertion 1]  
- [ ] [Inspection assertion 2]
```

2.4 CodeModel (code-model.js)

Responsible for FR-12 through FR-18. This is the runtime interface to the h-DSL instance.

Storage Layout

```
.puffin/
└── cre/
    ├── schema.json      # h-DSL schema definition
    ├── instance.json    # h-DSL instance (the code model)
    └── memo.json        # Cached navigation results
```

Key Methods

```
class CodeModel {
    /**
     * Load schema and instance from disk.
     * @param {string} projectRoot
     */
    async load(projectRoot) { /* ... */ }

    /**
     * Persist current state to disk.
     */
    async save() { /* ... */ }

    /**
     * PEEK: Get summary of an artifact without full load.
     * @param {string} path - Artifact path or identifier
     * @returns {ArtifactSummary}
     */
    peek(path) { /* ... */ }

    /**
     * FOCUS: Load full artifact details.
     * @param {string} path
     * @returns {Artifact}
     */
    focus(path) { /* ... */ }

    /**
     * TRACE: Follow dependency chain.
     * @param {string} from - Starting artifact
     * @param {Object} opts - { direction, kind, depth }
     * @returns {Artifact[]}
     */
    trace(from, opts) { /* ... */ }

    /**
     * FILTER: Find artifacts matching criteria.
     * @param {Object} criteria - { tags, kind, pattern }
     * @returns {Artifact[]}
     */
    filter(criteria) { /* ... */ }

    /**
     * Apply incremental updates to the instance.
     * @param {ModelDelta[]} deltas
     */
    async update(deltas) { /* ... */ }

    /**
     * Query artifacts relevant to a task description.
     * Implements ORIENT → FILTER → EXPLORE from the h-DSL protocol.
     * @param {string} task - Natural language task description
     * @returns {Artifact[]}
     */
    async queryForTask(task) { /* ... */ }
}
```

Model Delta Format

```
/**  
 * @typedef {Object} ModelDelta  
 * @property {string} op - add|update|remove  
 * @property {string} type - artifact|dependency|flow  
 * @property {string} path - Target path/identifier  
 * @property {Object} data - New or updated data  
 */
```

2.5 SchemaManager (schema-manager.js)

Responsible for schema evolution (FR-14, FR-15, FR-16).

```
class SchemaManager {  
    /**  
     * Load the current h-DSL schema.  
     * @returns {Schema}  
     */  
    async load() { /* ... */ }  
  
    /**  
     * Extend the schema with a new element type.  
     * Validates h-M3 annotation is present.  
     * @param {SchemaExtension} extension  
     */  
    async extend(extension) { /* ... */ }  
  
    /**  
     * Validate that an instance conforms to the schema.  
     * @param {Object} instance  
     * @returns {ValidationResult}  
     */  
    validate(instance) { /* ... */ }  
}
```

2.6 Introspector (introspector.js)

Responsible for FR-13 (post-implementation code model updates).

```
class Introspector {  
    /**  
     * Analyze code changes from a completed implementation.  
     * @param {string} branch - Branch that was implemented  
     * @param {string} baseBranch - Branch it was based on  
     * @returns {ModelDelta[]}  
     */  
    async analyzeChanges(branch, baseBranch) { /* ... */ }  
  
    /**  
     * Extract artifacts from source files.  
     * Parses files to identify modules, classes, functions, exports.  
     * @param {string[]} filePaths - Changed files  
     * @returns {Artifact[]}  
     */  
    async extractArtifacts(filePaths) { /* ... */ }  
  
    /**  
     * Infer intent and relationships from code structure.  
     * Uses AI to generate PROSE descriptions.  
     * @param {Artifact[]} artifacts  
     * @returns {ModelDelta[]}  
     */  
    async inferIntent(artifacts) { /* ... */ }  
}
```

2.7 AssertionGenerator (assertion-generator.js)

Responsible for FR-04.

```
class AssertionGenerator {  
    /**  
     * Generate inspection assertions for a plan item.  
     * @param {PlanItem} planItem  
     * @param {CodeModel} codeModel  
     * @returns {InspectionAssertion[]}  
     */  
    async generate(planItem, codeModel) { /* ... */ }  
  
    /**  
     * Run assertions against the current codebase.  
     * @param {InspectionAssertion[]} assertions  
     * @param {string} projectRoot  
     * @returns {AssertionResult[]}  
     */  
    async verify(assertions, projectRoot) { /* ... */ }  
}
```

3. h-DSL Schema Design

3.1 Base Schema (schema.json)

The schema defines what element types exist in the code model and how they map to h-M3.

```
{
  "version": "1.0.0",
  "m3Version": "2.0",
  "elementTypes": {
    "module": {
      "m3Type": "SLOT",
      "fields": {
        "path": { "m3Type": "TERM", "required": true },
        "kind": { "m3Type": "TERM", "required": true, "enum": ["module", "file", "config"] },
        "summary": { "m3Type": "PROSE", "required": true },
        "intent": { "m3Type": "PROSE", "required": false },
        "exports": { "m3Type": "TERM", "array": true },
        "tags": { "m3Type": "TERM", "array": true },
        "size": { "m3Type": "TERM", "required": false }
      }
    },
    "function": {
      "m3Type": "SLOT",
      "fields": {
        "path": { "m3Type": "TERM", "required": true },
        "signature": { "m3Type": "TERM", "required": false },
        "summary": { "m3Type": "PROSE", "required": true },
        "intent": { "m3Type": "PROSE", "required": false },
        "behavior": {
          "pre": { "m3Type": "PROSE" },
          "post": { "m3Type": "PROSE" },
          "err": { "m3Type": "PROSE" }
        },
        "tags": { "m3Type": "TERM", "array": true }
      }
    },
    "dependency": {
      "m3Type": "RELATION",
      "fields": {
        "from": { "m3Type": "TERM", "required": true },
        "to": { "m3Type": "TERM", "required": true },
        "kind": { "m3Type": "TERM", "required": true, "enum": ["imports", "calls", "extends", "implements", "configures", "tests"] },
        "weight": { "m3Type": "TERM", "enum": ["critical", "normal", "weak"] },
        "intent": { "m3Type": "PROSE", "required": false }
      }
    },
    "flow": {
      "m3Type": "SLOT",
      "fields": {
        "name": { "m3Type": "TERM", "required": true },
        "summary": { "m3Type": "PROSE", "required": true },
        "steps": {
          "array": true,
          "fields": {
            "order": { "m3Type": "TERM" },
            "artifact": { "m3Type": "TERM" },
            "intent": { "m3Type": "PROSE" }
          }
        },
        "tags": { "m3Type": "TERM", "array": true }
      }
    }
  },
  "extensionLog": []
}
```

3.2 Instance Format (`instance.json`)

```
{  
  "schemaVersion": "1.0.0",  
  "lastUpdated": "2026-01-31T00:00:00Z",  
  "artifacts": {  
    "src/main/plugin-loader.js": {  
      "type": "module",  
      "path": "src/main/plugin-loader.js",  
      "kind": "file",  
      "summary": "Discovers and loads plugins from the plugins directory",  
      "exports": ["loadPlugins", "validatePlugin"],  
      "tags": ["core", "plugins", "lifecycle"]  
    }  
  },  
  "dependencies": [  
    {  
      "from": "src/main/plugin-loader.js",  
      "to": "src/main/plugin-loader.js:PluginManager",  
      "kind": "calls",  
      "weight": "critical"  
    }  
  ],  
  "flows": {  
    "plugin-activation": {  
      "name": "plugin-activation",  
      "summary": "Plugin discovery through activation",  
      "steps": [  
        { "order": 1, "artifact": "PluginLoader.loadPlugins", "intent": "Scan plugins directory" },  
        { "order": 2, "artifact": "Plugin.activate", "intent": "Run plugin activate hook" }  
      ],  
      "tags": ["lifecycle", "startup"]  
    }  
  }  
}
```

3.3 Schema Evolution

When the Introspector encounters a concept not expressible in the current schema:

1. The Introspector identifies the gap (e.g., a "middleware" concept with ordering semantics).
2. It proposes a `SchemaExtension` with the new element type and its h-M3 annotation.
3. The `SchemaManager` validates the extension and appends it to `schema.json`.
4. Existing instance data remains valid (additive-only changes).

```
/**  
 * @typedef {Object} SchemaExtension  
 * @property {string} name - New element type name  
 * @property {string} m3Type - h-M3 primitive it maps to  
 * @property {Object} fields - Field definitions  
 * @property {string} rationale - Why this extension is needed  
 */
```

Extensions are logged in `schema.json.extensionLog` for auditability.

4. Database Schema

New tables added to the existing SQLite database.


```
-- Plans: one per sprint
CREATE TABLE plans (
    id TEXT PRIMARY KEY,
    sprint_id TEXT NOT NULL UNIQUE,
    status TEXT NOT NULL DEFAULT 'idle',
    version INTEGER NOT NULL DEFAULT 1,
    content TEXT NOT NULL, -- JSON serialized plan items
    questions TEXT, -- JSON array of pending questions
    created_at TEXT NOT NULL,
    updated_at TEXT NOT NULL,
    FOREIGN KEY (sprint_id) REFERENCES sprints(id)
);

-- Ready-to-Implement Specifications: many per user story
CREATE TABLE ris (
    id TEXT PRIMARY KEY,
    user_story_id TEXT NOT NULL,
    plan_id TEXT NOT NULL,
    sprint_id TEXT NOT NULL,
    branch TEXT NOT NULL,
    content TEXT NOT NULL, -- Markdown RIS content
    status TEXT NOT NULL DEFAULT 'generated',
    code_model_version TEXT, -- Hash of code model at generation time
    created_at TEXT NOT NULL,
    FOREIGN KEY (user_story_id) REFERENCES user_stories(id),
    FOREIGN KEY (plan_id) REFERENCES plans(id),
    FOREIGN KEY (sprint_id) REFERENCES sprints(id)
);


-- Inspection assertions: many per plan item
CREATE TABLE inspection_assertions (
    id TEXT PRIMARY KEY,
    plan_id TEXT NOT NULL,
    user_story_id TEXT NOT NULL,
    type TEXT NOT NULL, -- file_exists, function_exists, export_exists, pattern_match
    target TEXT NOT NULL, -- What to check
    description TEXT NOT NULL,
    result TEXT, -- pass|fail|pending
    verified_at TEXT,
    FOREIGN KEY (plan_id) REFERENCES plans(id),
    FOREIGN KEY (user_story_id) REFERENCES user_stories(id)
);
```

5. Sequence Diagrams


5.1 Plan Generation Flow



5.2 RIS Generation Flow



5.3 Post-Implementation Code Model Update



6. AI Integration Points

The CRE delegates reasoning to 3CLI (Claude) at specific points. These are not arbitrary AI calls — each maps to an h-M3 interpreter operation.

Operation	h-M3 Op	Where Used	Input	Output
Analyze ambiguities	GROUND	PlanGenerator.analyzeSprint	User stories	Clarifying questions
Generate plan	FOLLOW	PlanGenerator.generatePlan	Stories + answers + model	Ordered plan
Refine plan	FOLLOW	PlanGenerator.refinePlan	Plan + feedback	Updated plan
Generate assertions	DERIVE	AssertionGenerator.generate	Plan items	Testable assertions
Generate RIS	FOLLOW	RISGenerator.generateRIS	Story + plan + model	RIS markdown

Operation	h-M3 Op	Where Used	Input	Output
Infer intent	GROUND	Introspector.inferIntent	Source code	PROSE descriptions
Assess relevance	JUDGE	CodeModel.queryForTask	Task + artifacts	Ranked artifacts
Identify schema gaps	DERIVE	Introspector.analyzeChanges	New concepts	Schema extensions

Each AI call is structured as a prompt with: 1. **System context:** The relevant h-DL fragments and conventions. 2. **Task:** The specific operation (e.g., "generate an implementation plan"). 3. **Constraints:** Format requirements, length limits, required fields.

7. Error Handling Strategy

Per NFR-02, the CRE must not crash Puffin.

7.1 Error Categories

Category	Example	Recovery
AI Failure	3CLI returns malformed output	Retry once, then surface error to user
Schema Violation	Instance doesn't match schema	Log warning, skip invalid element, continue
Storage Failure	Can't write JSON/SQLite	Surface error to user, keep in-memory state
Model Inconsistency	Code model references deleted file	Mark artifact as stale, continue

Category	Example	Recovery
Plan Conflict	User story dependencies form cycle	Report cycle to user during plan generation

7.2 Error Wrapping

All IPC handlers wrap their logic:

```
ipcMain.handle('cre:generate-plan', async (event, { sprintId }) => {
  try {
    return { success: true, data: await planGenerator.analyzeSprint(sprintId) };
  } catch (error) {
    console.error('[CRE] Plan generation failed:', error);
    return { success: false, error: error.message };
  }
});
```

8. Configuration

CRE configuration lives in Puffin's existing config system.

```
{
  "cre": {
    "codeModelPath": ".puffin/cre",
    "maxPlanIterations": 5,
    "risMaxLength": 5000,
    "introspection": {
      "autoAfterMerge": true,
      "excludePatterns": ["node_modules/**", "dist/**", ".git/**"]
    },
    "schema": {
      "allowAutoExtension": true,
      "extensionApprovalRequired": false
    }
  }
}
```

9. Dependencies and Integration

9.1 Internal Dependencies

Dependency	Usage
SQLite database	Plan and RIS persistence

Dependency	Usage
Memory Plugin	Read branch memory, conversation history for context
Outcome Lifecycle Plugin	Read outcome states to inform planning
3CLI Runner	Execute AI prompts for reasoning operations
Git integration	Diff analysis for introspection

9.2 Integration with Orchestration Layer

The orchestration layer drives the CRE through IPC calls in this sequence:

1. User starts sprint → `cre:generate-plan`
 2. User answers questions → `cre:refine-plan` (repeat as needed)
 3. User approves → `cre:approve-plan`
 4. For each story in order → `cre:generate-ris`
 5. 3CLI implements RIS
 6. After implementation + review + fixes → `cre:update-model`
 7. Repeat from step 4 for next story
-

10. Open Questions

#	Question	Impact
OQ-1	Should the code model be versioned (git-tracked) or treated as derived state?	Affects recovery strategy and storage approach
OQ-2	What is the maximum acceptable latency for plan generation?	May require streaming/progressive output
OQ-3	Should RIS generation be parallelizable across independent user stories?	Affects orchestration and model locking

#	Question	Impact
OQ-4	How should the CRE handle codebase changes made outside of Puffin?	May need periodic full re-introspection
OQ-5	Should schema extensions require user approval?	Trade-off between autonomy and control